

Feature	Single-row Subquery	Correlated Subquery
Execution	Runs once	Runs once per outer row
Dependence	Independent of outer query	References outer query columns
Example Use	Compare salary > company avg	Compare salary > dept avg
Performance	Faster (1 execution)	Slower (many executions possible)

👉 Quick interview tip:

- If you see the inner subquery referring to an **outer alias (like f.department_id)** → it's **correlated**.
 - If not → it's a **single-row (or multi-row) subquery**.
-

Feature	Multi-Row Subquery	Multi-Column Subquery
Definition	Subquery that returns one column with multiple rows	Subquery that returns multiple columns (tuples) with multiple rows
Operators Used	<code>IN, ANY, ALL</code>	Tuple comparison with <code>IN</code> or <code>EXISTS</code>
Outer Query Comparison	Compares a single column from outer query against a set of values from inner query	Compares multiple columns together as a single tuple
Example	<pre>sql SELECT first_name, fare FROM titanic WHERE fare > ANY (SELECT fare FROM titanic WHERE pclass = 1); → Compares one column (fare) to many fares.</pre>	<pre>sql SELECT * FROM titanic e WHERE (e.pclass, e.embark_town) IN (SELECT t.pclass, t.embark_town FROM titanic t WHERE survived = 1); → Compares two columns (pclass, embark_town) as a pair.</pre>

Quick Identifier	Inner query: <code>SELECT col1 ...</code>	Inner query: <code>SELECT col1, col2 ...</code> Outer query: <code>(col1, col2)</code> in tuple form
Use Case	<ul style="list-style-type: none"> - Find values in a single column that match multiple rows (e.g., fares, ages, IDs). 	<ul style="list-style-type: none"> - Match on combinations of values (e.g., class + embark_town, dept_id + job_id).

Quick Differentiation Trick

- **Look at the subquery SELECT:**
 - If it has **1 column** → Multi-Row.
 - If it has **2+ columns** → Multi-Column.
 - Outer query will use either:
 - `col IN (SELECT col ...)` → Multi-Row.
 - `(col1, col2) IN (SELECT col1, col2 ...)` → Multi-Column.
-

Feature	NOT IN	NOT EXISTS
How it works	Compares a value to a list returned by the subquery.	Checks whether the subquery returns any row at all .
Syntax	<code>sql WHERE col NOT IN (SELECT col FROM table ...)</code>	<code>sql WHERE NOT EXISTS (SELECT 1 FROM table WHERE ...)</code>
NULL behavior	⚠️ If the subquery returns even one NULL , the whole <code>NOT IN</code> returns no rows	✓ Ignores NULLs automatically. Works safely even if subquery has NULLs.

(because comparisons with NULL = unknown).

Performance	Usually fine for small sets; can be slower if subquery list is big.	Often faster in large tables, since it stops at the first match .
Use case	Use when you're sure the subquery column doesn't contain NULLs (e.g., primary keys).	Use when NULLs are possible, or when you want to be 100% safe.
Example	Passengers not on decks where survivors were: <pre>sql SELECT * FROM titanic e WHERE e.deck NOT IN (SELECT deck FROM titanic WHERE survived=1 AND deck IS NOT NULL);</pre>	Same query safe with NOT EXISTS : <pre>sql SELECT * FROM titanic e WHERE NOT EXISTS (SELECT 1 FROM titanic t WHERE t.deck = e.deck AND t.survived=1);</pre>

Quick rules to remember:

- **Prefer NOT EXISTS** → safer (handles NULLs correctly).
- **NOT IN** → only when you're 100% sure the subquery column has **no NULL values** (like IDs, primary keys).
- Many interviewers like to test this: “*What happens if the subquery returns NULL?*” — the answer: **NOT IN** fails, **NOT EXISTS** works fine.

Feature	Correlated Subquery	Multi-Row Subquery
Definition	Inner subquery depends on a column from the outer query. Runs once for each row of the outer query.	Inner subquery returns multiple rows (a set) that the outer query compares against. Runs once total , independently.

Execution	Repeated execution per outer row.	Single execution, returns a list/set.
Operators used	Normal comparison operators (<code>=, >, <</code>) with correlation to outer alias.	<code>IN, ANY, ALL</code> are commonly used to handle multiple values.
Example (Titanic dataset)	<p>Passengers who paid more than the average fare of their embark_town:</p> <pre>sql SELECT f.passenger_no, f.fare FROM titanic f WHERE f.fare > (SELECT AVG(t.fare) FROM titanic t WHERE t.embark_town = f.embark_town);</pre> <p>→ inner query uses <code>f.embark_town</code> from the outer query → correlated.</p>	<p>Passengers in the same pclass as survivors:</p> <pre>sql SELECT passenger_no, first_name, pclass, survived FROM titanic WHERE pclass IN (SELECT DISTINCT pclass FROM titanic WHERE survived = 1);</pre> <p>→ inner query returns a list of classes → multi-row.</p>
When to use	When you need a calculation or condition that changes per row of the outer query (e.g., compare against that row's department avg).	When you need to check if a value is in a set of results (e.g., department is in the list of departments that meet a condition).
Performance	Can be slower on big data (runs many times). Often rewritten with JOINs or window functions.	Usually faster, since inner query runs once and returns a set.

Quick trick to differentiate

- If the **inner query references outer query columns** → it's **correlated**.
 - If the **inner query stands alone and returns multiple rows** → it's **multi-row**.
-

Feature	IN	EXISTS
Definition	Compares a value from the outer query to a list of values returned by the subquery.	Tests whether the subquery returns at least one row . Doesn't care about values.

Execution	Subquery runs first → returns a set → outer query checks if value is in that set.	Outer query runs, then for each row, subquery checks if a matching row exists (correlated).
Return Type	Subquery must return one column (can be many rows).	Subquery can return any column(s) , SQL only checks row existence.
NULL behavior	⚠️ If subquery contains even one NULL , NOT IN may return no rows at all (because <code>col NOT IN (... , NULL)</code> is always unknown).	✓ NOT EXISTS safely handles NULLs (it only checks existence of rows, not values).
Performance	Better for small static lists (e.g., <code>WHERE col IN (1,2,3)</code>).	Often better for large, correlated queries (stops after finding the first match).
Example (Titanic dataset)	Passengers in the same pclass as survivors: <code>sql SELECT passenger_no, pclass FROM titanic WHERE pclass IN (SELECT DISTINCT pclass FROM titanic WHERE survived = 1);</code>	Passengers in the same pclass as survivors: <code>sql SELECT e.passenger_no, e.pclass FROM titanic e WHERE EXISTS (SELECT 1 FROM titanic t WHERE t.pclass = e.pclass AND t.survived = 1);</code>

Quick rules for interviews

- ✓ Use **IN** when comparing to a **list/set of values**.
 - ✓ Use **EXISTS** when you just want to check if a matching row exists.
 - ✓ Prefer **EXISTS** over **NOT IN** if NULLs might appear.
-

Updating rooms vs bookings – why derived table was needed

Updating **rooms** – no problem

```
UPDATE rooms r
```

```
SET r.status = 'High Demand'
```

```

WHERE EXISTS (
    SELECT b.roomnumber, COUNT(*) AS c
    FROM bookings b
    WHERE b.roomnumber = r.roomnumber
    GROUP BY b.roomnumber
    HAVING c > 5
);

```

1.

- Here we **update rooms** but **read from bookings** in the subquery.
- Different tables → **MySQL is happy**, no error.

Updating **bookings** – direct subquery causes error 1093

```

UPDATE bookings
SET bookingstatus = 'High demand'
WHERE totalamount > (
    SELECT AVG(totalamount) FROM bookings -- ✗ same table
);

```

2.

- MySQL rule: **you can't update a table and read from the same table in a subquery** → error 1093.

Fix: use a **derived table** (temp result)

```

UPDATE bookings c
SET c.bookingstatus = 'High demand'
WHERE EXISTS (
    SELECT roomnumber
    FROM (

```

```

SELECT roomnumber, COUNT(*) AS c
FROM bookings b
GROUP BY roomnumber
HAVING c > 5
) AS Raj           -- ✓ derived table
WHERE Raj.roomnumber = c.roomnumber -- ✓ correlation
);

```

3.

- Inner query on `bookings` is materialized as `Raj` (a temp table).
 - Outer UPDATE reads from `Raj`, not directly from `bookings` → **no 1093 error**.
 - Correlation `Raj.roomnumber = c.roomnumber` ensures **only bookings of high-demand rooms** get updated.
-

 **Memory line:**

When `UPDATE` and subquery use the **same table**, MySQL needs a **derived table (or JOIN)** to avoid error 1093.

Different table (like `rooms` reading `bookings`) → no problem.